



PageRank & HITS

CE642: Social and Economic Networks
Maryam Ramezani
Sharif University of Technology
maryam.ramezani@sharif.edu

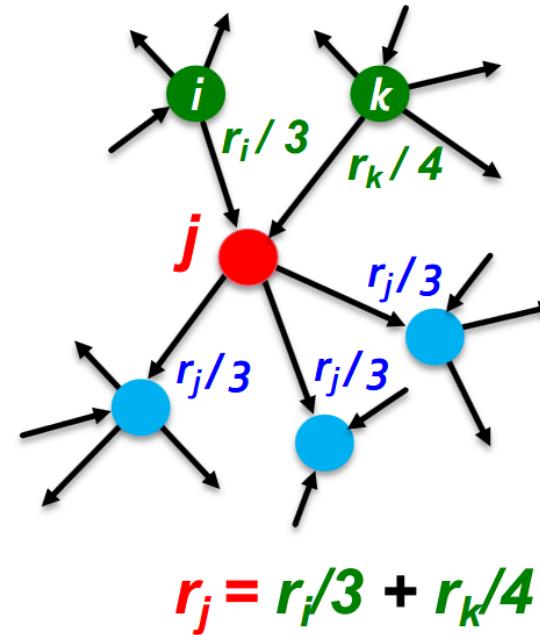


01

PageRank

PageRank

- A “vote” from an important page is worth more:
 - Each link’s vote is proportional to the **importance** of its source page
 - If page i with importance r_i has d_i out-links, each link gets r_i / d_i votes
 - Page j ’s own importance r_j is the sum of the votes on its in-links



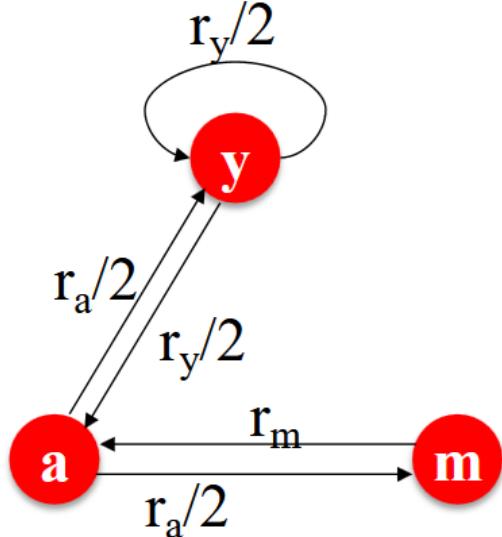
PageRank: Flow View

- A page is important if it is pointed to by other important pages
- Define “rank” r_j for node j

$$r_j = \sum_{i \rightarrow j} \frac{r_i}{d_i}$$

d_i ... out-degree of node i

You might wonder: Let's just use Gaussian elimination to solve this system of linear equations. Bad idea!



“Flow” equations:

$$r_y = r_y/2 + r_a/2$$

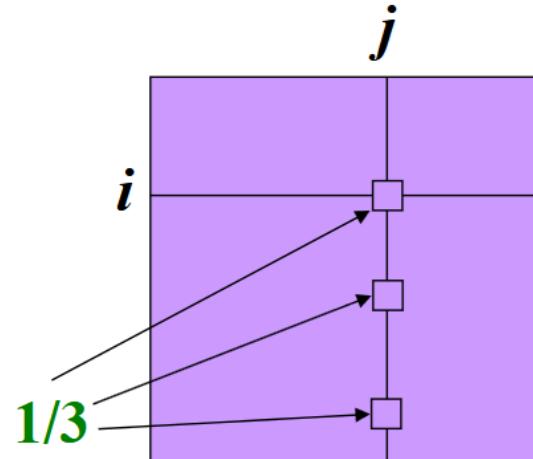
$$r_a = r_y/2 + r_m$$

$$r_m = r_a/2$$

PageRank: Matrix View

- **Stochastic adjacency matrix M**

- Let page j have d_j out-links
 - If $j \rightarrow i$, then $M_{ij} = \frac{1}{d_j}$
 - M is a **column stochastic matrix**
 - Columns sum to 1



- **Rank vector r :** An entry per page

M

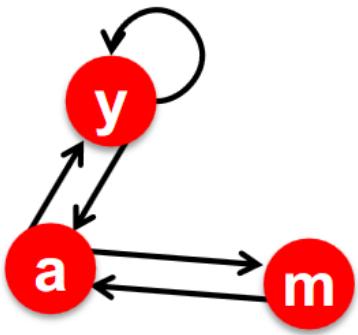
- r_i is the importance score of page i
 - $\sum_i r_i = 1$

- **The flow equations can be written**

$$\mathbf{r} = \mathbf{M} \cdot \mathbf{r}$$

$$r_j = \sum_{i \rightarrow j} \frac{r_i}{d_i}$$

PageRank Example



	r_y	r_a	r_m
r_y	$\frac{1}{2}$	$\frac{1}{2}$	0
r_a	$\frac{1}{2}$	0	1
r_m	0	$\frac{1}{2}$	0

$$r_y = r_y/2 + r_a/2$$

$$r_a = r_y/2 + r_m$$

$$r_m = r_a/2$$

$$\begin{bmatrix} r_y \\ r_a \\ r_m \end{bmatrix} = \begin{bmatrix} \frac{1}{2} & \frac{1}{2} & 0 \\ \frac{1}{2} & 0 & 1 \\ 0 & \frac{1}{2} & 0 \end{bmatrix} \begin{bmatrix} r_y \\ r_a \\ r_m \end{bmatrix}$$

r M r

PageRank: Matrix View

$$C_p(v_i) = \alpha \sum_{j=1}^n A_{j,i} \frac{C_p(v_j)}{d_j^{\text{out}}} + \beta$$

What if the degree is zero?

$$\begin{cases} d_j^{\text{out}} > 0 \\ D = \text{diag}(d_1^{\text{out}}, d_2^{\text{out}}, \dots, d_n^{\text{out}}) \end{cases} \rightarrow \mathbf{C}_p = \alpha A^T D^{-1} \mathbf{C}_p + \beta \mathbf{1}$$

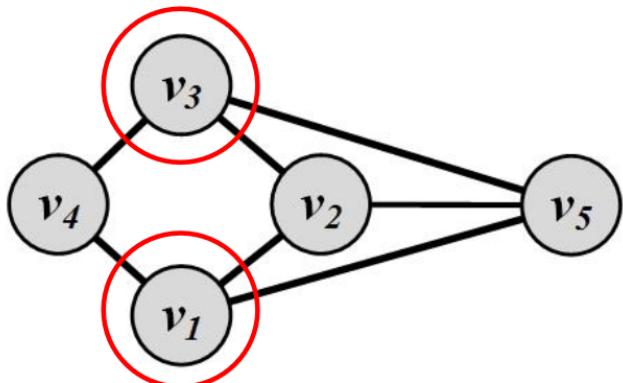


$$\mathbf{C}_p = \beta(\mathbf{I} - \alpha A^T D^{-1})^{-1} \cdot \mathbf{1}$$

Similar to Katz Centrality, in practice, $\alpha < 1/\lambda$, where λ is the largest eigenvalue of $A^T D^{-1}$. In undirected graphs, the largest eigenvalue of $A^T D^{-1}$ is $\lambda = 1$; therefore, $\alpha < 1$.

PageRank Example

- We assume $\alpha=0.95 < 1$ and $\beta = 0.1$



$$A = \begin{bmatrix} 0 & 1 & 0 & 1 & 1 \\ 1 & 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 1 & 1 \\ 1 & 0 & 1 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \end{bmatrix}$$

$$\mathbf{C}_p = \beta(\mathbf{I} - \alpha A^T D^{-1})^{-1} \cdot \mathbf{1} = \begin{bmatrix} 2.14 \\ 2.13 \\ 2.14 \\ 1.45 \\ 2.13 \end{bmatrix}$$

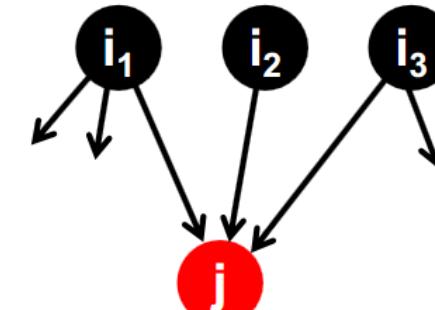
Connection to Random Walk

- **Imagine a random web surfer:**

- At any time t , surfer is on some page i
- At time $t + 1$, the surfer follows an out-link from i uniformly at random
- Ends up on some page j linked from i
- Process repeats indefinitely

- **Let:**

- $p(t)$... vector whose i^{th} coordinate is the prob. that the surfer is at page i at time t
- So, $p(t)$ is a probability distribution over pages



$$r_j = \sum_{i \rightarrow j} \frac{r_i}{d_{\text{out}}(i)}$$

Stationary Distribution

- Where is the surfer at time $t+1$?

- Follow a link uniformly at random

$$p(t + 1) = M \cdot p(t)$$

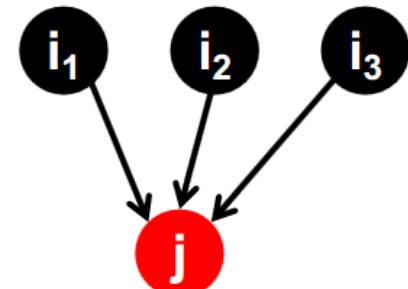
- Suppose the random walk reaches a state

$$p(t + 1) = M \cdot p(t) = p(t)$$

then $p(t)$ is **stationary distribution** of a random walk

- Our original rank vector r satisfies $r = M \cdot r$

- So, r is a stationary distribution for the random walk



$$p(t + 1) = M \cdot p(t)$$

How to solve PageRank?

- The flow equation:

$$1 \cdot \mathbf{r} = \mathbf{M} \cdot \mathbf{r}$$

$$\begin{matrix} \mathbf{r}_y \\ \mathbf{r}_a \\ \mathbf{r}_m \end{matrix} = \begin{matrix} \frac{1}{2} & \frac{1}{2} & 0 \\ \frac{1}{2} & 0 & 1 \\ 0 & \frac{1}{2} & 0 \end{matrix} \begin{matrix} \mathbf{r}_y \\ \mathbf{r}_a \\ \mathbf{r}_m \end{matrix}$$

$$\mathbf{r} \quad \mathbf{M} \quad \mathbf{r}$$

- So the **rank vector** \mathbf{r} is an **eigenvector** of the stochastic adj. matrix \mathbf{M} (with eigenvalue 1)
 - Starting from any vector \mathbf{u} , the limit $\mathbf{M}(\mathbf{M}(\dots \mathbf{M}(\mathbf{M} \mathbf{u})))$ is the **long-term distribution** of the surfers.
 - **PageRank** = Limiting distribution = **principal eigenvector** of M
 - **Note:** If \mathbf{r} is the limit of the product $\mathbf{M}\mathbf{M} \dots \mathbf{M}\mathbf{u}$, then \mathbf{r} satisfies the **flow equation** $1 \cdot \mathbf{r} = \mathbf{M}\mathbf{r}$
 - So \mathbf{r} is the **principal eigenvector** of M with eigenvalue 1
- **We can now efficiently solve for r !**
 - The method is called **Power iteration**

Power Iteration for PageRank

- Given a web graph with N nodes, where the nodes are pages and edges are hyperlinks
- Power iteration: a simple iterative scheme

- Initialize: $\mathbf{r}^0 = [1/N, \dots, 1/N]^T$

- Iterate: $\mathbf{r}^{(t+1)} = \mathbf{M} \cdot \mathbf{r}^t$

- Stop when $|\mathbf{r}^{(t+1)} - \mathbf{r}^t|_1 < \varepsilon$

$$r_j^{(t+1)} = \sum_{i \rightarrow j} \frac{r_i^{(t)}}{d_i}$$

d_i out-degree of node i

$|x|_1 = \sum_1^N |x_1|$ is the L₁ norm

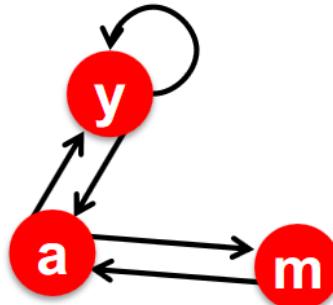
Can use any other vector norm, e.g., Euclidean

About 50 iterations is sufficient to estimate the limiting solution.

PageRank Example

■ Power Iteration:

- Set $r_j \leftarrow 1/N$
- 1: $r'_j \leftarrow \sum_{i \rightarrow j} \frac{r_i}{d_i}$
- 2: If $|r - r'| > \varepsilon$:
 - $r \leftarrow r'$
- 3: go to 1



	y	a	m
y	$\frac{1}{2}$	$\frac{1}{2}$	0
a	$\frac{1}{2}$	0	1
m	0	$\frac{1}{2}$	0

$$\begin{aligned}r_y &= r_y/2 + r_a/2 \\r_a &= r_y/2 + rm \\r_m &= r_a/2\end{aligned}$$

■ Example:

$$\begin{bmatrix} r_y \\ r_a \\ r_m \end{bmatrix} = \begin{bmatrix} 1/3 \\ 1/3 \\ 1/3 \end{bmatrix}, \begin{bmatrix} 1/3 \\ 3/6 \\ 1/6 \end{bmatrix}, \begin{bmatrix} 5/12 \\ 1/3 \\ 3/12 \end{bmatrix}, \begin{bmatrix} 9/24 \\ 11/24 \\ 1/6 \end{bmatrix}, \dots, \begin{bmatrix} 6/15 \\ 6/15 \\ 3/15 \end{bmatrix}$$

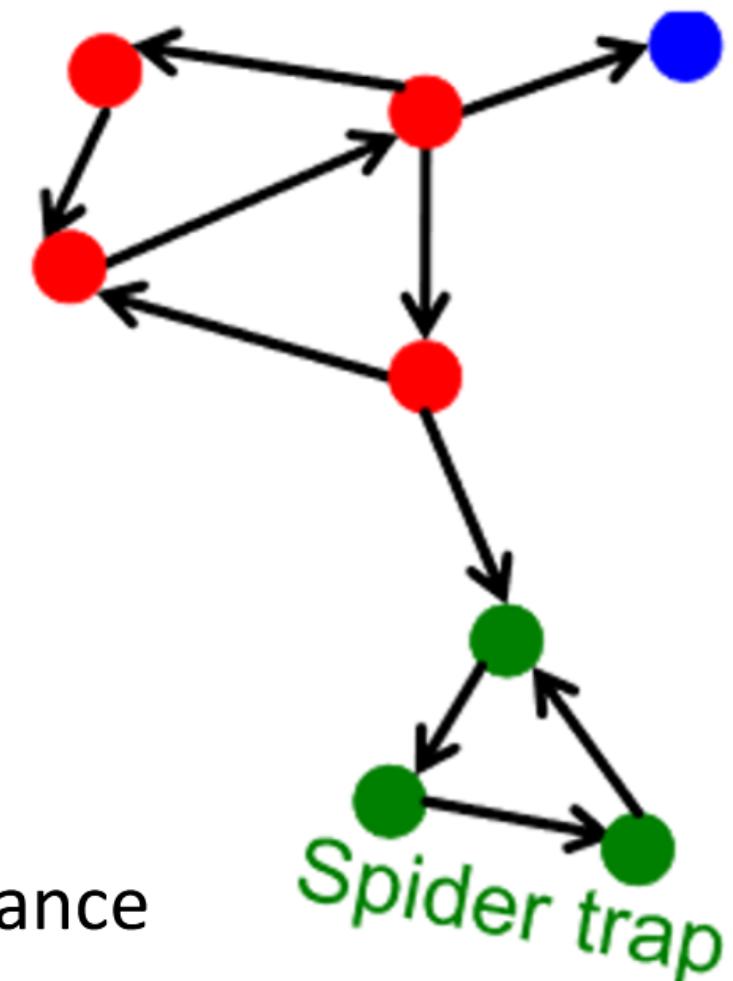
Iteration 0, 1, 2, ...

PageRank Problems

Two problems:

- (1) Some pages are **dead ends** (have no out-links)
 - Such pages cause importance to “leak out”

- (2) **Spider traps**
(all out-links are within the group)
 - Eventually spider traps absorb all importance



“Dead End” Problem



$$r_j^{(t+1)} = \sum_{i \rightarrow j} \frac{r_i^{(t)}}{d_i}$$

■ Example:

$$\begin{array}{lcl} r_a & = & 1 \quad 0 \quad 0 \quad 0 \\ r_b & & 0 \quad 1 \quad 0 \quad 0 \end{array}$$

Iteration 0, 1, 2, ...

“Dead End” Problem



$$r_j^{(t+1)} = \sum_{i \rightarrow j} \frac{r_i^{(t)}}{d_i}$$

■ Example:

$$\begin{array}{lcl} r_a & = & 1 \quad 0 \quad 0 \quad 0 \\ r_b & & 0 \quad 1 \quad 0 \quad 0 \end{array}$$

Iteration 0, 1, 2, ...

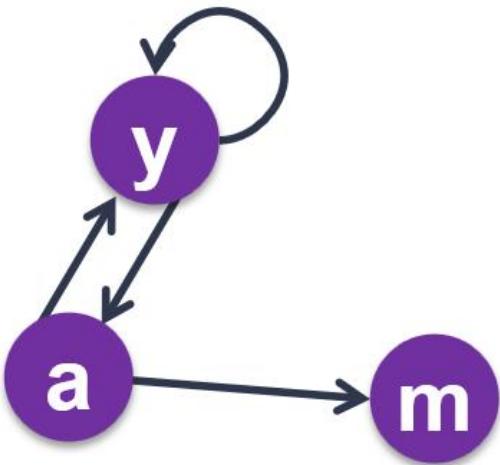
“Dead End” Problem

Power Iteration:

Set $r_j = 1$

$$r_j = \sum_{i \rightarrow j} \frac{r_i}{d_i}$$

And iterate



$$\begin{bmatrix} r_y \\ r_a \\ r_m \end{bmatrix} = \begin{matrix} 1/3 & 2/6 & 3/12 & 5/24 & \dots & 0 \\ 1/3 & 1/6 & 2/12 & 3/24 & \dots & 0 \\ 1/3 & 1/6 & 1/12 & 2/24 & \dots & 0 \end{matrix}$$

Iteration 0, 1, 2, ...

Here the PageRank “leaks” out since the matrix is not stochastic.

	y	a	m
y	$\frac{1}{2}$	$\frac{1}{2}$	0
a	$\frac{1}{2}$	0	0
m	0	$\frac{1}{2}$	0

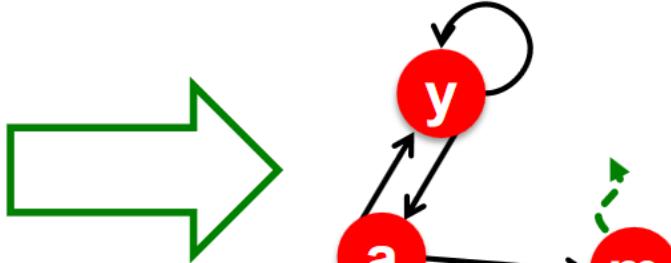
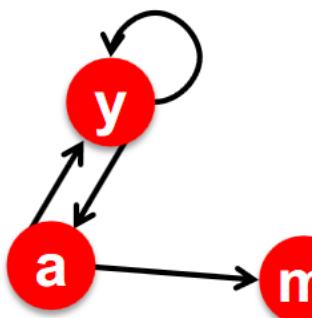
$$r_y = r_y/2 + r_a/2$$

$$r_a = r_y/2$$

$$r_m = r_a/2$$

Solution to “Dead End” Problem

- **Teleports:** Follow random teleport links with total probability **1.0** from dead-ends
 - Adjust matrix accordingly

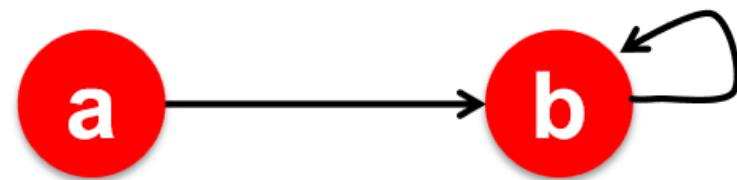


	y	a	m
y	$\frac{1}{2}$	$\frac{1}{2}$	0
a	$\frac{1}{2}$	0	0
m	0	$\frac{1}{2}$	0

	y	a	m
y	$\frac{1}{2}$	$\frac{1}{2}$	$\frac{1}{3}$
a	$\frac{1}{2}$	0	$\frac{1}{3}$
m	0	$\frac{1}{2}$	$\frac{1}{3}$

“Spider Trap” Problem

- The “Spider trap” problem:



$$r_j^{(t+1)} = \sum_{i \rightarrow j} \frac{r_i^{(t)}}{d_i}$$

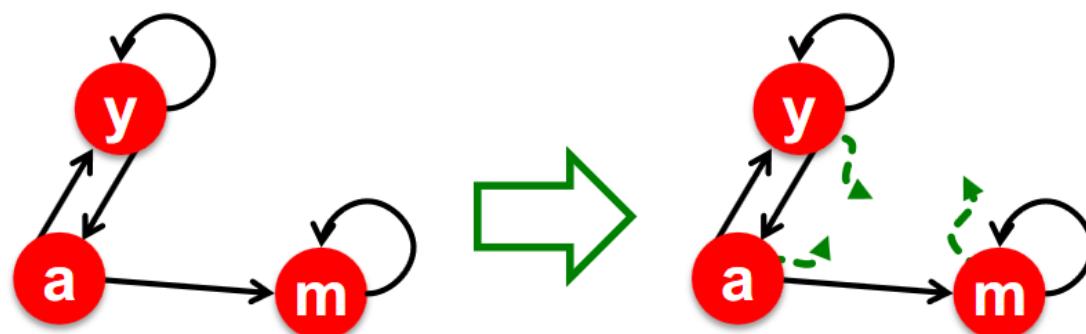
- Example:

Iteration: 0, 1, 2, 3...

r_a	=	1		0		0		0
r_b		0		1		1		1

Solution to “Spider Trap” Problem

- Solution for spider traps: At each time step, the random surfer has two options
 - With prob. β , follow a link at random
 - With prob. $1-\beta$, jump to a random page
 - Common values for β are in the range 0.8 to 0.9
- Surfer will teleport out of spider trap within a few time steps



Why Teleports Solve the Problem?

Why are dead-ends and spider traps a problem
and why do teleports solve the problem?

- **Spider-traps** are not a problem, but with traps
PageRank scores are **not** what we want
 - **Solution:** Never get stuck in a spider trap by teleporting out of it in a finite number of steps
- **Dead-ends** are a problem
 - The matrix is not column stochastic so our initial assumptions are not met
 - **Solution:** Make matrix column stochastic by always teleporting when there is nowhere else to go

The Google Matrix

- Google's solution that does it all:

At each step, random surfer has two options:

- With probability β , follow a link at random
- With probability $1-\beta$, jump to some random page

- **PageRank equation** [Brin-Page, 98]

$$r_j = \sum_{i \rightarrow j} \beta \frac{r_i}{d_i} + (1 - \beta) \frac{1}{N}$$

d_i ... out-degree
of node i

This formulation assumes that M has no dead ends. We can either preprocess matrix M to remove all dead ends or explicitly follow random teleport links with probability 1.0 from dead-ends.

The Google Matrix

- **PageRank equation** [Brin-Page, '98]

$$r_j = \sum_{i \rightarrow j} \beta \frac{r_i}{d_i} + (1 - \beta) \frac{1}{N}$$

- **The Google Matrix G :**

$[1/N]_{N \times N}$...N by N matrix
where all entries are 1/N

$$G = \beta M + (1 - \beta) \left[\frac{1}{N} \right]_{N \times N}$$

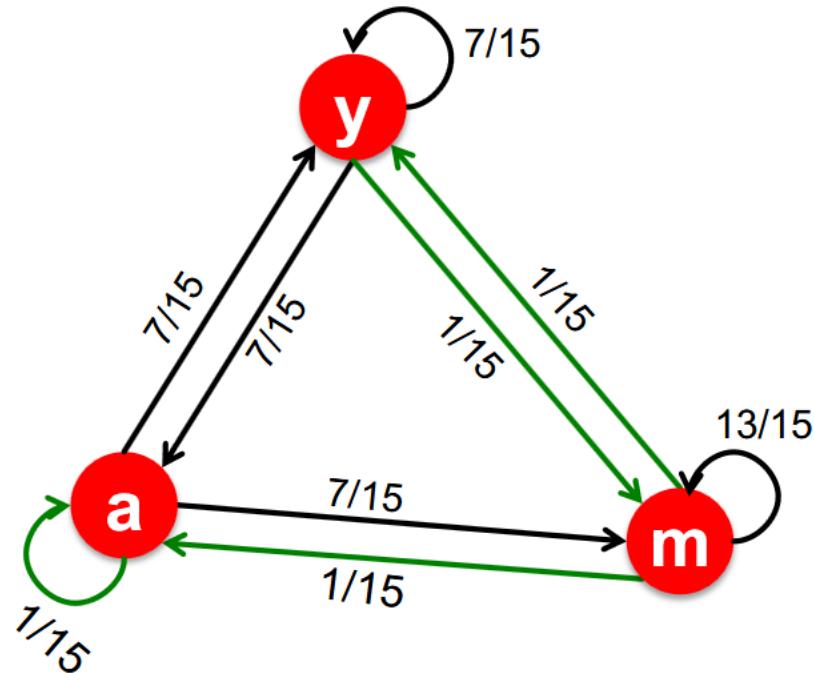
- **We have a recursive problem:** $r = G \cdot r$

And the Power method still works!

- **What is β ?**

- In practice $\beta = 0.8, 0.9$ (make 5 steps on avg., jump)

Random Teleports ($\beta = 0.8$)



$$\text{M} \quad [1/N]_{N \times N}$$
$$0.8 \quad \begin{matrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 0 \\ 0 & 1/2 & 1 \end{matrix} + 0.2 \quad \begin{matrix} 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \end{matrix}$$
$$G \quad \begin{matrix} y & 7/15 & 7/15 & 1/15 \\ a & 7/15 & 1/15 & 1/15 \\ m & 1/15 & 7/15 & 13/15 \end{matrix}$$

y	1/3	0.33	0.24	0.26	...	7/33
a	=	1/3	0.20	0.20	0.18	5/33
m		1/3	0.46	0.52	0.56	21/33

Conclusion

- PageRank solves for $r = Gr$ and can be efficiently computed by power iteration of the stochastic adjacency matrix (G)
- Adding random uniform teleportation solves issues of dead-ends and spider-traps

PageRank Problems

- **Measures generic popularity of a page**
 - Biased against topic-specific authorities
 - **Solution:** Topic-Specific PageRank (next)
- **Uses a single measure of importance**
 - Other models e.g., [hubs-and-authorities](#)
 - **Solution:** Hubs-and-Authorities (next)
- **Susceptible to Link spam**
 - Artificial link topographies created in order to boost page rank
 - **Solution:** TrustRank (next)



02

Topic-Specific PageRank

Topic-Specific PageRank

- Instead of generic popularity, can we measure popularity within a topic?
- Goal: Evaluate Web pages not just according to their popularity, but by how close they are to a particular topic, e.g. “sports” or “history.”
- Allows search queries to be answered based on interests of the user
 - Example: Query “Trojan” wants different pages depending on whether you are interested in sports or history.

Topic-Specific PageRank

- Assume each walker has a small probability of “teleporting” at any step
- **Teleport can go to:**
 - Any page with equal probability
 - To avoid dead-end and spider-trap problems
 - A topic-specific set of “relevant” pages (teleport set)
 - For topic-sensitive PageRank.
- **Idea: Bias the random walk**
 - When walked teleports, she pick a page from a set S
 - S contains only pages that are relevant to the topic
 - E.g., Open Directory (DMOZ) pages for a given topic
 - For each teleport set S , we get a different vector r_S

Topic-Specific PageRank

- Let:

- $$\begin{aligned} A_{ij} &= \beta M_{ij} + (1-\beta) / |S| && \text{if } i \in S \\ &= \beta M_{ij} && \text{otherwise} \end{aligned}$$

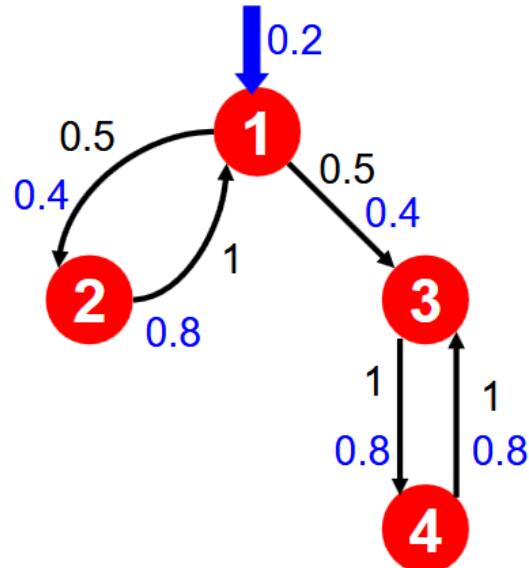
- A is stochastic!
- We have weighted all pages in the teleport set S equally

- Could also assign different weights to pages!

- Compute as for regular PageRank:

- Multiply by M, then add a vector
 - Maintains sparseness

Topic-Specific PageRank Example



Suppose $S = \{1\}$, $\beta = 0.8$

Node	Iteration			
	0	1	2...	stable
1	1.0	0.2	0.52	0.294
2	0	0.4	0.08	0.118
3	0	0.4	0.08	0.327
4	0	0	0.32	0.261

Note how we initialize the PageRank vector differently from the unbiased PageRank case.



03

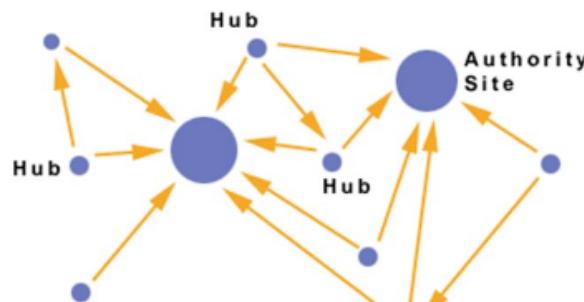
HITS (Hypertext-Induced Topic Selection)

Hubs and Authorities

Interesting pages fall into two classes:

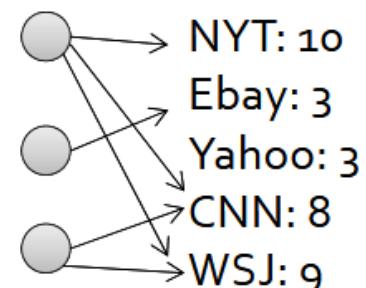
1. **Authorities** are pages containing useful information

- Newspaper home pages
- Course home pages
- Home pages of auto manufacturers



2. **Hubs** are pages that link to authorities

- List of newspapers
- Course bulletin
- List of US auto manufacturers



Hubs and Authorities

■ Hubs and Authorities

Each page has 2 scores:

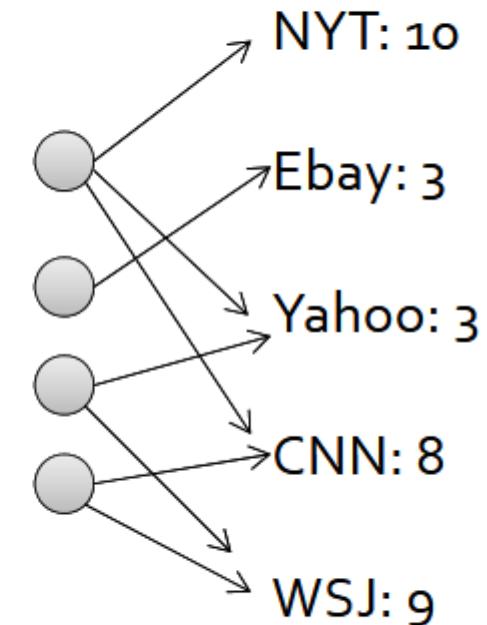
- **Quality as an expert (hub):**

- Total sum of votes of pages pointed to

- **Quality as an content (authority):**

- Total sum of votes of experts

- Principle of repeated improvement



Hubs and Authorities

- A good hub links to many good authorities
- A good authority is linked from many good hubs
- Model using two scores for each node:
 - Hub score and Authority score
 - Represented as vectors h and a

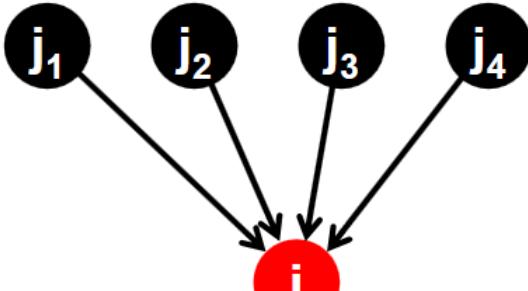
Hubs and Authorities

- Each page i has 2 scores:

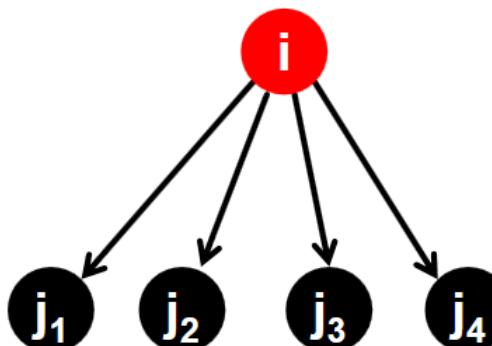
- Authority score: a_i
- Hub score: h_i

HITS algorithm:

- Initialize: $a_j = 1, h_i = 1$
- Then keep iterating:
 - $\forall i$: Authority: $a_i = \sum_{j \rightarrow i} h_j$
 - $\forall i$: Hub: $h_i = \sum_{i \rightarrow j} a_j$
 - $\forall i$: normalize: $\sum_j a_j = 1, \sum_j h_j = 1$



$$a_i = \sum_{j \rightarrow i} h_j$$



$$h_i = \sum_{i \rightarrow j} a_j$$

Transition Matrix A

- HITS converges to a single stable point
- Slightly change the notation:

- Vector $a = (a_1, \dots, a_n)$, $h = (h_1, \dots, h_n)$

- Adjacency matrix ($n \times n$): $A_{ij} = 1$ if $i \rightarrow j$

- Then:

$$h_i = \sum_{i \rightarrow j} a_j \Leftrightarrow h_i = \sum_j A_{ij} a_j$$

- So: $h = A a$
- And likewise: $a = A^T h$

Hubs and Authorities Equations

- The **hub** score of page i is proportional to the sum of the **authority** scores of the pages it links to: $h = \lambda A a$
 - Constant λ is a scale factor, $\lambda=1/\sum h_i$
- The **authority** score of page i is proportional to the sum of the **hub** scores of the pages it is linked from: $a = \mu A^T h$
 - Constant μ is scale factor, $\mu=1/\sum a_i$

Iterative Algorithm

- The HITS algorithm:
 - Initialize \mathbf{h} , \mathbf{a} to all 1's
 - Repeat:
 - $\mathbf{h} = \mathbf{A} \mathbf{a}$
 - Scale \mathbf{h} so that its sums to 1.0
 - $\mathbf{a} = \mathbf{A}^T \mathbf{h}$
 - Scale \mathbf{a} so that its sums to 1.0
 - Until \mathbf{h} , \mathbf{a} converge (i.e., change very little)

Hubs and Authorities Equations

- HITS algorithm in new notation:

- Set: $a = h = I^n$

- Repeat:

- $h = A a, \quad a = A^T h$

- Normalize

- Then: $a = A^T \underbrace{(A a)}_{\text{new } h}$
 $\qquad\qquad\qquad \underbrace{\text{new } a}_{\text{new } a}$

- Thus, in $2k$ steps:

$$a = (A^T A)^k a$$

$$h = (A A^T)^k h$$

a is being updated (in 2 steps):

$$A^T(A a) = (A^T A) a$$

h is updated (in 2 steps):

$$A(A^T h) = (A A^T) h$$

Repeated matrix powering

Hubs and Authorities Equations

- $h = \lambda A a$ $\lambda = 1/\sum h_i$
 - $a = \mu A^T h$ $\mu = 1/\sum a_i$
 - $h = \lambda \mu A A^T h$
 - $a = \lambda \mu A^T A a$
 - Under reasonable assumptions about A , the HITS iterative algorithm converges to vectors h^* and a^* :
 - h^* is the principal eigenvector of matrix $A A^T$
 - a^* is the principal eigenvector of matrix $A^T A$

Conclusion

- PageRank and HITS are two solutions to the same problem:
 - What is the value of an in-link from u to v ?
 - In the PageRank model, the value of the link depends on the links into u
 - In the HITS model, it depends on the value of the other links out of u

04

TrustRank

Idea

- **Basic principle: Approximate isolation**
 - It is rare for a “good” page to point to a “bad” (spam) page
- Sample a set of “seed pages” from the web
- Have an **oracle** (human) identify the good pages and the spam pages in the seed set
 - **Expensive task**, so we must make seed set as small as possible

Idea

- Call the subset of seed pages that are identified as “good” the “trusted pages”
- Perform a topic-sensitive PageRank with teleport set = trusted pages.
 - Propagate trust through links:
 - Each page gets a trust value between 0 and 1
- Use a threshold value and mark all pages below the trust threshold as spam

Simple Model

- Set trust of each trusted page to 1
- Suppose trust of page p is t_p
 - Set of out-links o_p
- For each $q \in o_p$, p confers the trust:
 - $\beta t_p / |o_p|$ for $0 < \beta < 1$
- **Trust is additive**
 - Trust of p is the sum of the trust conferred on p by all its in-linked pages
- **Note similarity to Topic-Specific PageRank**
 - Within a scaling factor, TrustRank = PageRank with trusted pages as teleport set



Any Question?