

# Introduction to Graph Neural Networks

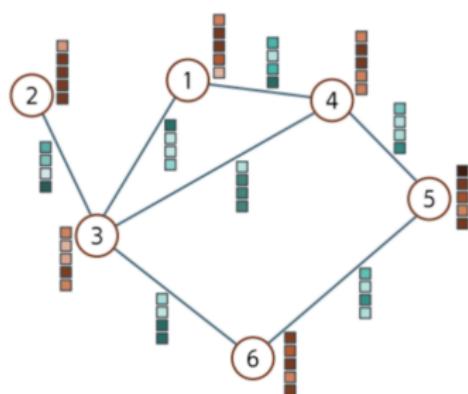
Mahdi Mastani

27 May 2025

# Graph Representation and Encoding

A graph consists of:

- **Graph Structure**
- **Node Embeddings**
- **Edge Embeddings**
- **Degree Matrix**



Adjacency  
matrix,  $\mathbf{A}$   
 $N \times N$

	1	2	3	4	5	6
1	0	1	1	1	0	0
2	0	0	1	0	0	0
3	1	0	0	1	1	1
4	1	0	1	0	1	0
5	0	0	1	1	0	1
6	0	0	1	0	1	0

Node  
data,  $\mathbf{X}$   
 $D \times N$

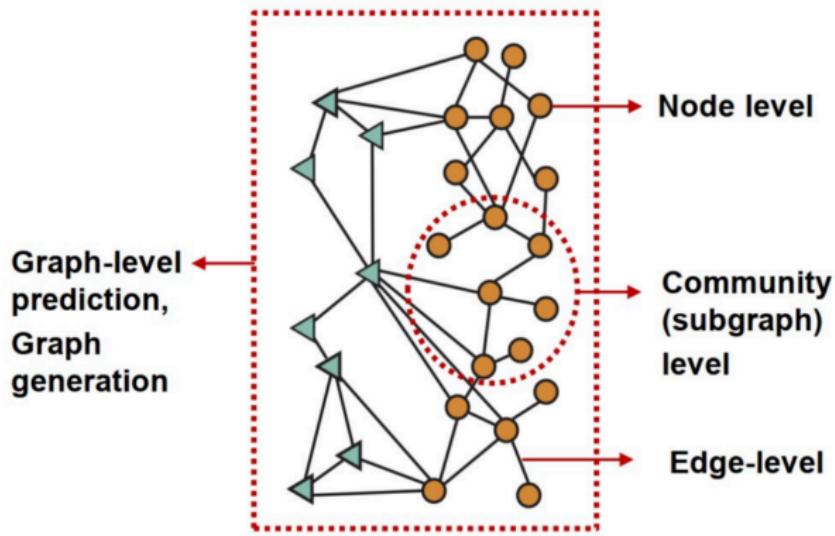
1	2	3	4	5	6
0.5	0.5	0.5	0.5	0.5	0.5
0.5	0.5	0.5	0.5	0.5	0.5
0.5	0.5	0.5	0.5	0.5	0.5
0.5	0.5	0.5	0.5	0.5	0.5
0.5	0.5	0.5	0.5	0.5	0.5
0.5	0.5	0.5	0.5	0.5	0.5

Edge  
data,  $\mathbf{E}$   
 $D_E \times E$

1	1	2	3	3	4	5
3	4	3	4	6	5	6
0.5	0.5	0.5	0.5	0.5	0.5	0.5
0.5	0.5	0.5	0.5	0.5	0.5	0.5
0.5	0.5	0.5	0.5	0.5	0.5	0.5
0.5	0.5	0.5	0.5	0.5	0.5	0.5
0.5	0.5	0.5	0.5	0.5	0.5	0.5

# Common Tasks for Graphs

- Node Classification
- Link Prediction
- Graph Classification
- Community Detection
- Graph Generation



# Key Properties of Graph Neural Networks

- **Generalization:** The ability to apply learned models to graphs of different sizes and topologies.
- **Scalability:** The architecture should be efficient enough to handle large graphs with millions of nodes and edges.
- **Permutation equivariance:** The model should produce the same output regardless of the ordering of the nodes and edges in the input graph.

# Permutation Equivariance in Graphs

- **Permutation matrix:** A permutation matrix  $P \in \{0, 1\}^{n \times n}$  is a binary square matrix with exactly one entry of 1 in each row and column. It represents a reordering of elements.

$$P = \begin{bmatrix} 0 & 0 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix}$$

- When position  $(i, j)$  of the permutation matrix is set to **one**, it indicates that node  $i$  will become node  $j$  after the permutation.

# Indexing and Permutation Effects

- Changing node indexing in a graph requires transforming the data accordingly.
- Pre-multiplying by  $P$  reorders the **rows** (used for node features).
- Post-multiplying by  $P^T$  reorders the **columns** (used for graph structure).
- The operations to map between indexings:

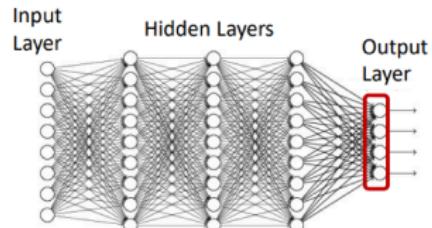
$$X' = PX, \quad A' = PAP^T$$

- **Conclusion:** Any graph processing model should remain invariant to these permutations:

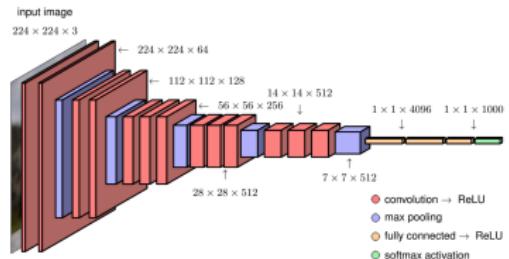
$$P\hat{y}(X, A) = \hat{y}(X', A')$$

# Neural Networks

- **MLP:**



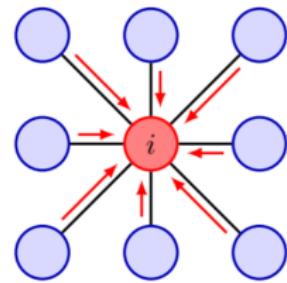
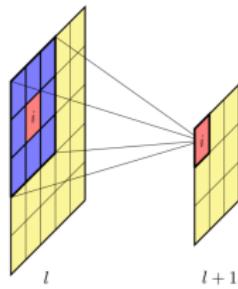
- **CNN:**



# From CNNs to GNNs

## How CNNs Work

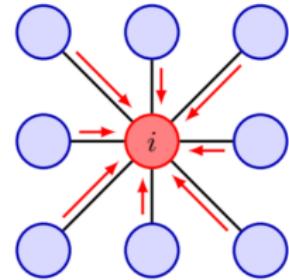
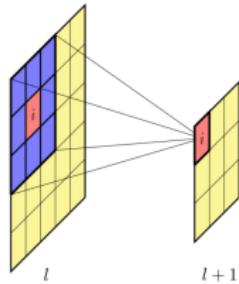
- CNNs operate on grid-structured data (like images).
- Use local filters (kernels) to scan spatially arranged data.
- Employ weight sharing and local connectivity to capture local patterns.



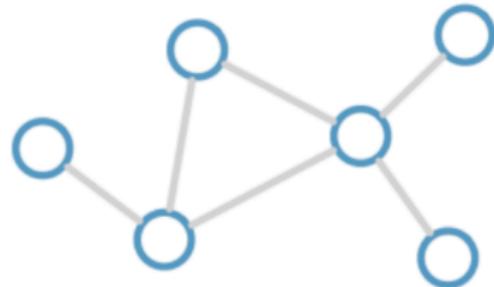
$$\mathbf{h}_i^{(l+1)} = \sigma \left( \sum_j \mathbf{w}_j^{(l)} \mathbf{h}_j^{(l)} \right)$$

# Why CNNs Are Not Suitable for Graphs

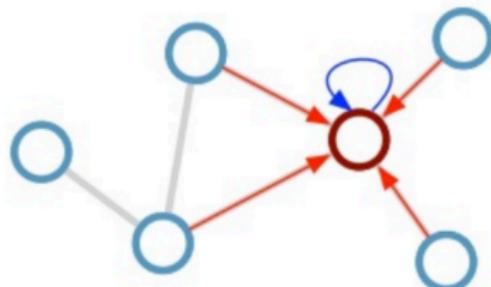
- Graphs are non-Euclidean: no fixed node order or grid structure.
- Nodes may have varying numbers of neighbors.



# GNN Message Passing - Neighborhood Aggregation



Undirected Graph



Update Rule for Node  
Embedding

- Step 1: Aggregate neighbors
- Step 2: Add self-loop

**Update rule:**  $\mathbf{h}_i^{(l+1)} = \sigma \left( \mathbf{h}_i^{(l)} \mathbf{W}_0^{(l)} + \sum_{j \in \mathcal{N}_i} \frac{1}{c_{ij}} \mathbf{h}_j^{(l)} \mathbf{W}_1^{(l)} \right)$

# Simple Message-Passing Neural Network

---

**Algorithm 1** Simple message-passing neural network

---

**Require:** Undirected graph  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

Initial node embeddings  $\{\mathbf{h}_n^{(0)} = \mathbf{x}_n\}$

Aggregate( $\cdot$ ) function

Update( $\cdot, \cdot$ ) function

**Ensure:** Final node embeddings  $\{\mathbf{h}_n^{(L)}\}$

```
1: // Iterative message-passing
2: for  $l \in \{0, \dots, L - 1\}$  do
3:    $\mathbf{z}_n^{(l)} \leftarrow \text{Aggregate}\left(\left\{\mathbf{h}_m^{(l)} : m \in \mathcal{N}(n)\right\}\right)$ 
4:    $\mathbf{h}_n^{(l+1)} \leftarrow \text{Update}\left(\mathbf{h}_n^{(l)}, \mathbf{z}_n^{(l)}\right)$ 
5: end for
6: return  $\{\mathbf{h}_n^{(L)}\}$ 
```

---

# Aggregator and Update Functions in GNNs

## Aggregator Function:

- Aggregator must be **permutation invariant**.
- Options:
  - **Sum**: Adds up neighbor features; sensitive to node degree.
  - **Mean**: Computes the average of neighbor features.
  - **Max**: Captures the most prominent signal per feature dimension.

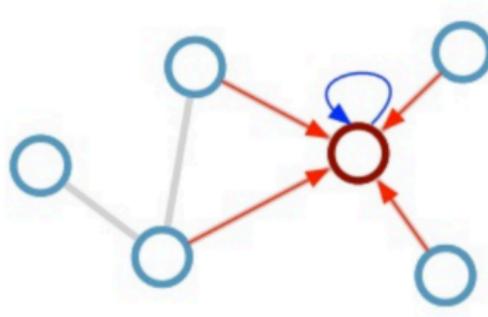
## Update Function:

- Update function should preserve or enhance node representations.
- Typically a neural network (e.g., MLP or linear layer).
- Can include residual connections or batch normalization.

# Building GCN Step-by-Step

## Step 1: Neighborhood Aggregation

**Update rule:**  $\mathbf{h}_i^{(l+1)} = \sigma \left( \mathbf{h}_i^{(l)} \mathbf{W}_0^{(l)} + \sum_{j \in \mathcal{N}_i} \frac{1}{c_{ij}} \mathbf{h}_j^{(l)} \mathbf{W}_1^{(l)} \right)$

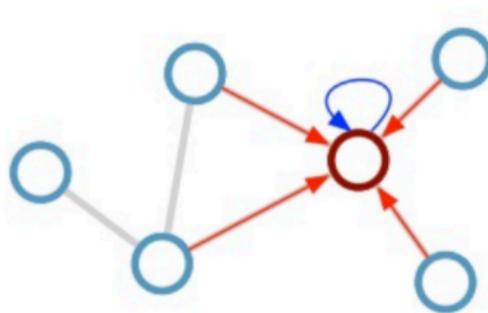


# Building GCN Step-by-Step

if we set  $\mathbf{W}_0^l = \mathbf{W}_1^l = \mathbf{W}^l$  (Shared weight matrix):

$$\mathbf{h}_i^{(l+1)} = \sigma \left( \mathbf{h}_i^{(l)} \mathbf{W}^{(l)} + \sum_{j \in \mathcal{N}_i} \frac{1}{c_{ij}} \mathbf{h}_j^{(l)} \mathbf{W}^{(l)} \right)$$

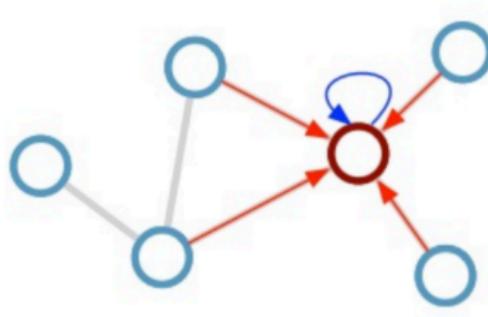
- how about  $c_{ij}$  ?



# Building GCN Step-by-Step

use Kipf normalization  $c_{ij} = \sqrt{d_i d_j}$

$$\mathbf{h}_i^{(l+1)} = \sigma \left( \mathbf{h}_i^{(l)} \mathbf{W}^{(l)} + \sum_{j \in \mathcal{N}_i} \frac{1}{\sqrt{d_i d_j}} \mathbf{h}_j^{(l)} \mathbf{W}^{(l)} \right)$$



# Translate to Graph Input

$$\mathbf{h}_i^{(l+1)} = \sigma \left( \mathbf{h}_i^{(l)} \mathbf{W}^{(l)} + \sum_{j \in \mathcal{N}_i} \frac{1}{\sqrt{d_i d_j}} \mathbf{h}_j^{(l)} \mathbf{W}^{(l)} \right)$$

Matrix form:

$$\mathbf{H}^{(l+1)} = \sigma \left( \mathbf{H}^{(l)} \mathbf{W}^{(l)} + \mathbf{D}^{-1/2} \mathbf{A} \mathbf{D}^{-1/2} \mathbf{H}^{(l)} \mathbf{W}^{(l)} \right)$$

- If set  $\hat{\mathbf{A}} = \mathbf{I} + \mathbf{D}^{-1/2} \mathbf{A} \mathbf{D}^{-1/2}$  we have:

$$\mathbf{H}^{(l+1)} = \sigma \left( \hat{\mathbf{A}} \mathbf{H}^{(l)} \mathbf{W}^{(l)} \right)$$

# L-layer GCN

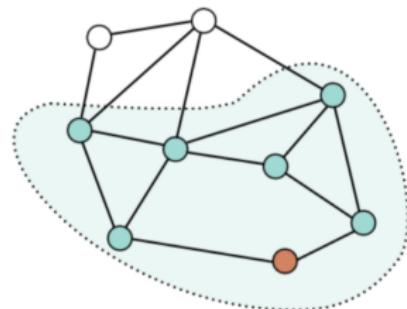
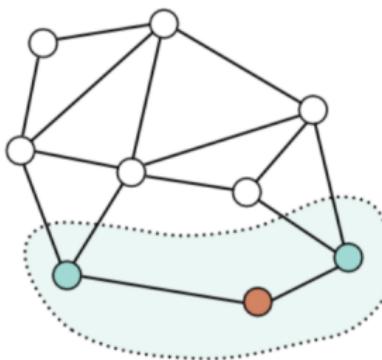
L-layer Graph convolutional networks (GCNs):

$$\mathbf{H}^{(1)} = \mathbf{F}(\mathbf{X}, \mathbf{A}, \mathbf{W}^{(1)}) = \sigma(\hat{\mathbf{A}}\mathbf{X}\mathbf{W}^{(1)})$$

$$\mathbf{H}^{(2)} = \mathbf{F}(\mathbf{H}^{(1)}, \mathbf{A}, \mathbf{W}^{(2)}) = \sigma(\hat{\mathbf{A}}\mathbf{H}^{(1)}\mathbf{W}^{(2)})$$

$$\vdots = \vdots$$

$$\mathbf{H}^{(L)} = \mathbf{F}(\mathbf{H}^{(L-1)}, \mathbf{A}, \mathbf{W}^{(L)}) = \sigma(\hat{\mathbf{A}}\mathbf{H}^{(L-1)}\mathbf{W}^{(L)})$$



# From Fixed to Learnable Coefficients

- So far, we discussed using fixed normalization coefficients  $\frac{1}{c_{ij}}$ , such as:
  - Uniform (unweighted average)
  - Degree-based normalization (e.g.,  $\frac{1}{\sqrt{d_i d_j}}$ )
- However, these do not adapt based on node features or context.
- Can we make these coefficients **learnable** instead?

# Graph Attention Layer - Overview

**Goal:** Compute hidden representations for each node by attending over its neighbors using self-attention.

## Key Properties:

- Efficient and parallelizable across node-neighbor pairs.
- Supports nodes with varying degrees using adaptive neighbor weights.
- Input features:  $\mathbf{h} = \{\vec{h}_1, \vec{h}_2, \dots, \vec{h}_n\}, \quad \vec{h}_i \in \mathbb{R}^F$
- Shared linear transformation:  $\mathbf{W} \in \mathbb{R}^{F' \times F}$

# Self-Attention Mechanism in GAT

## Step 1: Linear Transformation

$$\vec{h}'_i = \mathbf{W}\vec{h}_i \quad \forall i \in \mathcal{V}$$

## Step 2: Compute Attention Coefficients

$$s_{ij} = a(\vec{h}'_i, \vec{h}'_j)$$

Where  $a : \mathbb{R}^{F'} \times \mathbb{R}^{F'} \rightarrow \mathbb{R}$

### Popular Choices for Attention Scoring:

- **Dot Product:**  $a(\vec{h}'_i, \vec{h}'_j) = (\vec{h}'_i)^\top \vec{h}'_j$
- **Additive:**

$$a(\vec{h}'_i, \vec{h}'_j) = \text{LeakyReLU}(\mathbf{a}^\top [\vec{h}'_i \| \vec{h}'_j])$$

# Attention Score Matrix $\mathbf{S}$

- The matrix  $\mathbf{S} \in \mathbb{R}^{n \times n}$  contains raw attention scores:  $s_{ij} = a(\vec{h}'_i, \vec{h}'_j)$
- These scores indicate the importance of node  $j$  to node  $i$  based on transformed features.

So:

$$\mathbf{H}_{new} = \sigma(\mathbf{S} \cdot \mathbf{H}')$$

## Why not apply $\mathbf{S}$ directly?

- $\mathbf{S}$  is a matrix of **unnormalized scores** — directly using it can lead to unstable and unbounded outputs.
- If we set  $s'_{ij} = \frac{\exp(s_{ij})}{\sum_{k=1}^n \exp(s_{ik})}$  will have:

$$\mathbf{H}_{new} = \sigma(\mathbf{S}' \cdot \mathbf{H}')$$

# Attention Score Matrix $\mathbf{S}$

- It does not respect the graph structure — it includes all node-to-node interactions unless masked.
- It may cause unrelated nodes to influence each other.
- Mask scores outside neighborhood  $\mathcal{N}_i$

$$\mathbf{s}_{ij}' = \frac{\exp(s_{ij})}{\sum_{k \in \mathcal{N}_i \cup i} \exp(s_{ik})}$$

- Resulting matrix  $\mathbf{S}$  is row-stochastic (i.e., values sum to 1 per row)

# Masked Attention in Matrix Form

## Final Attention Mechanism with Graph Structure:

- $M = A + I$
- Softmax over masked positions requires attention scores to be set to  $-\infty$  for excluded elements
- Zero values in  $M$  set to  $-\infty$
- Apply masking before softmax:

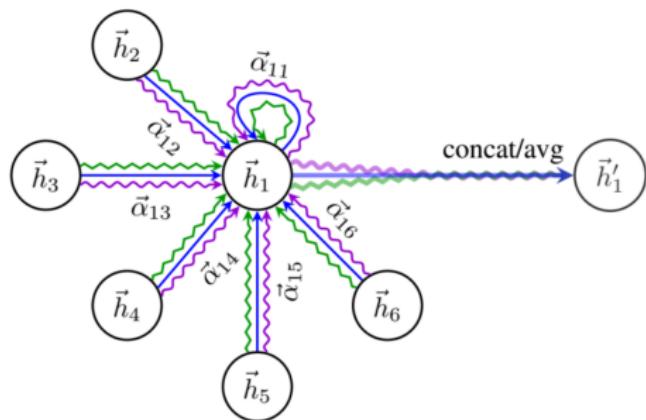
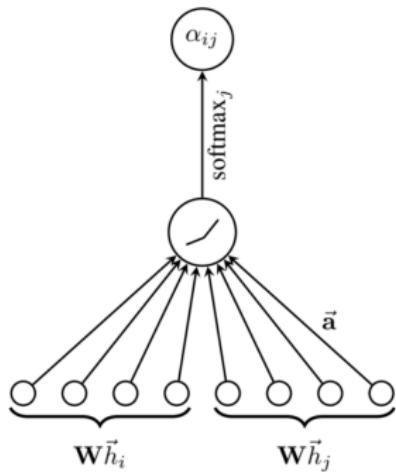
$$\tilde{S} = \text{softmax}(S \odot M)$$

- $\odot$ : element-wise multiplication (masking)

## Final Update Rule:

$$H_{\text{new}} = \sigma(\tilde{S}H')$$

# GAT



# Training GNNs: Supervised and Unsupervised Settings

## What if we don't have any labels? (Unsupervised Learning)

- Use node features and graph structure to learn useful representations.
- **One possible idea:** "Similar" nodes should have similar embeddings.

$$\min_W \mathcal{L} = \sum_{u,v} \text{CE}(y_{u,v}, \langle \vec{h}_v, \vec{h}_u \rangle)$$

- $y_{u,v} = 1$  if node  $u$  and  $v$  are **similar**
- $\langle \vec{h}_v, \vec{h}_u \rangle$ : **similarity of embeddings**
- **Node Similarity** can be:
  - edges
  - Random walk distance

# Supervised GNN Training: Node Classification

**Task:** Predict a label  $y_i$  for each node  $i \in \mathcal{V}$

**Approach:**

- Use GNN to compute node embeddings  $\vec{h}_i$
- Apply a softmax classifier on each embedding

**Loss Function: Cross-Entropy**

$$\mathcal{L} = - \sum_{i \in \mathcal{V}_{\text{labeled}}} \sum_{c=1}^C y_{ic} \log \hat{y}_{ic}$$

where  $\hat{y}_{ic} = \text{softmax}(\mathbf{W}\vec{h}_i)$

# Supervised GNN Training: Graph Classification

**Task:** Predict a label for the entire graph  $G$

**Approach:**

- Compute node embeddings  $\vec{h}_v$
- Aggregate (e.g., mean, sum, attention) to form graph embedding  $\vec{h}_G$
- Apply a classifier on  $\vec{h}_G$

**Loss Function: Cross-Entropy (for classification)**

$$\mathcal{L} = - \sum_{G \in \mathcal{D}} \sum_{c=1}^C y_{Gc} \log \hat{y}_{Gc}$$

where  $\hat{y}_{Gc} = \text{softmax}(\mathbf{W}\vec{h}_G)$

# Supervised GNN Training: Link Prediction

**Task:** Predict whether an edge exists between a node pair  $(u, v)$

**Approach:**

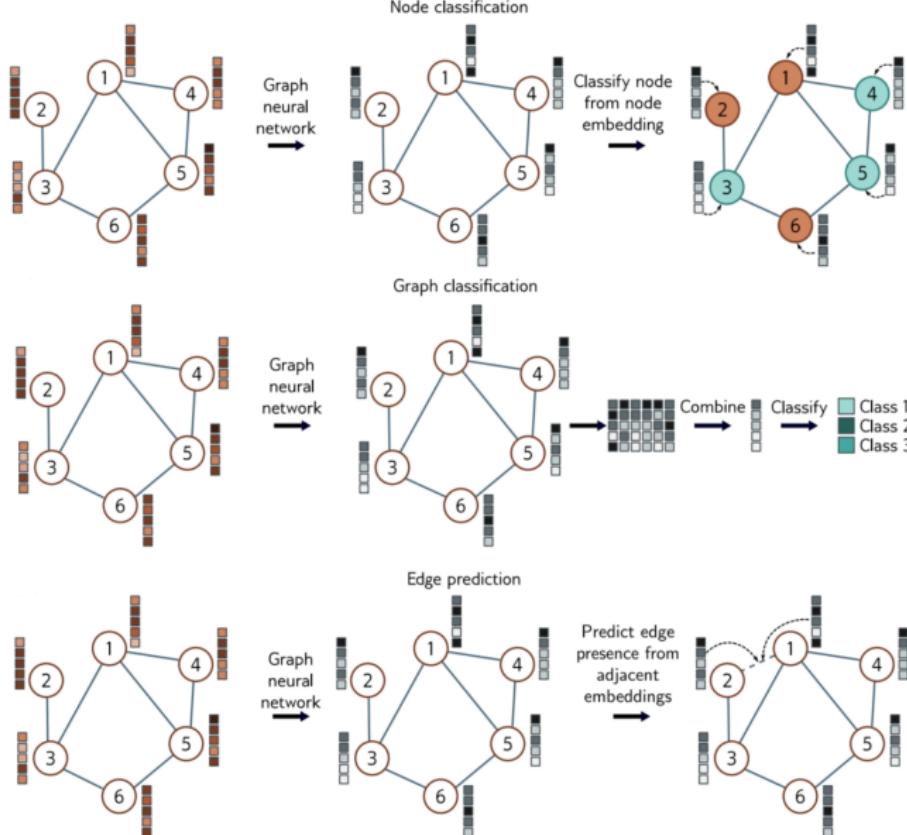
- Use GNN to compute embeddings  $\vec{h}_u, \vec{h}_v$
- Predict link score using dot product or MLP:

$$\hat{y}_{uv} = \sigma(\vec{h}_u^\top \vec{h}_v) \quad \text{or} \quad \text{MLP}([\vec{h}_u \| \vec{h}_v])$$

**Loss Function: Binary Cross-Entropy**

$$\mathcal{L} = - \sum_{(u,v)} y_{uv} \log \hat{y}_{uv} + (1 - y_{uv}) \log (1 - \hat{y}_{uv})$$

# Visualization of Tasks



# Semi-Supervised Learning

**Definition:** Learning from a dataset that contains both **labeled** and **unlabeled** examples.

**Occurs When:**

- Labels are expensive or time-consuming to obtain.
- Large amounts of raw (unlabeled) data are available.

**Objective:**

- Use unlabeled data to improve generalization.
- Learn representations that respect both labels and data structure.

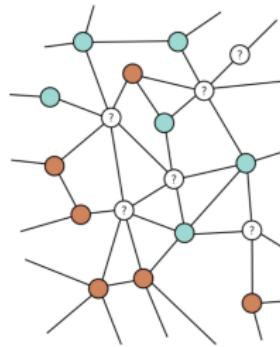
# Semi-Supervised Learning in Graphs

## Problem Setup:

- A graph  $G = (\mathcal{V}, \mathcal{E})$  with node features.
- Only a subset of nodes  $\mathcal{V}_L \subset \mathcal{V}$  are labeled.

## Key Idea:

- Use both graph structure and node features to propagate labels.
- Unlabeled nodes benefit from neighboring labeled information via message passing.



# Inductive vs. Transductive Learning in GNNs

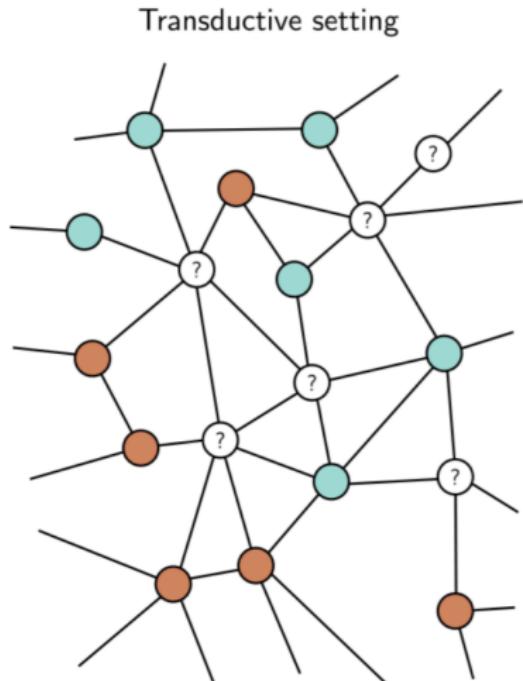
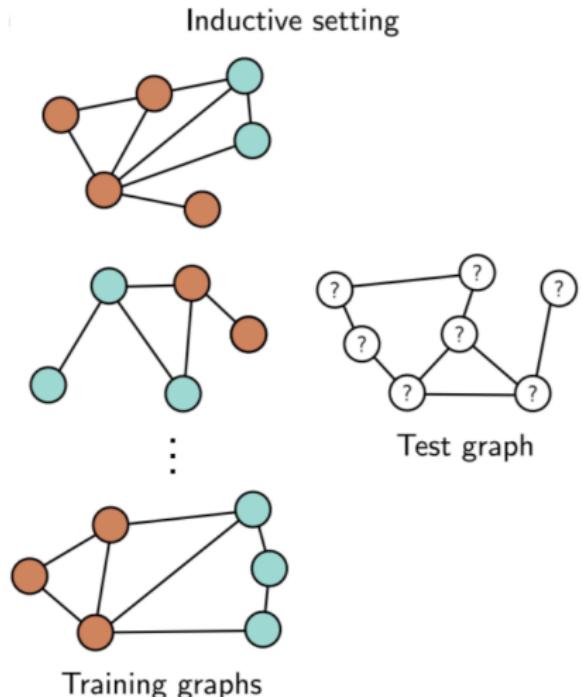
## Inductive Learning:

- Learns a general rule from labeled training data that maps inputs to outputs.
- Once trained, the model can be applied to **new, unseen data**.
- This is the default in most machine learning settings.

## Transductive Learning:

- Considers both labeled and unlabeled data **simultaneously** during training.
- Does not learn a reusable rule — instead directly infers labels for the current test nodes.
- Can exploit patterns in unlabeled data, but must be retrained if new data are added.

# Inductive vs. Transductive Learning in GNNs



# Training Large Graphs in Batches

**Challenge:** Large graphs may not fit into memory, making full-graph training impractical.

## Mini-batch Training Strategies:

- **Layer-wise sampling:** Sample fixed-size sets of neighbors per GNN layer.
- **Graph partitioning:** Cluster the original graph into disjoint subsets of nodes.

# Graph Partitioning

**Goal:** Break down a large graph into smaller, more manageable subgraphs for mini-batch training.

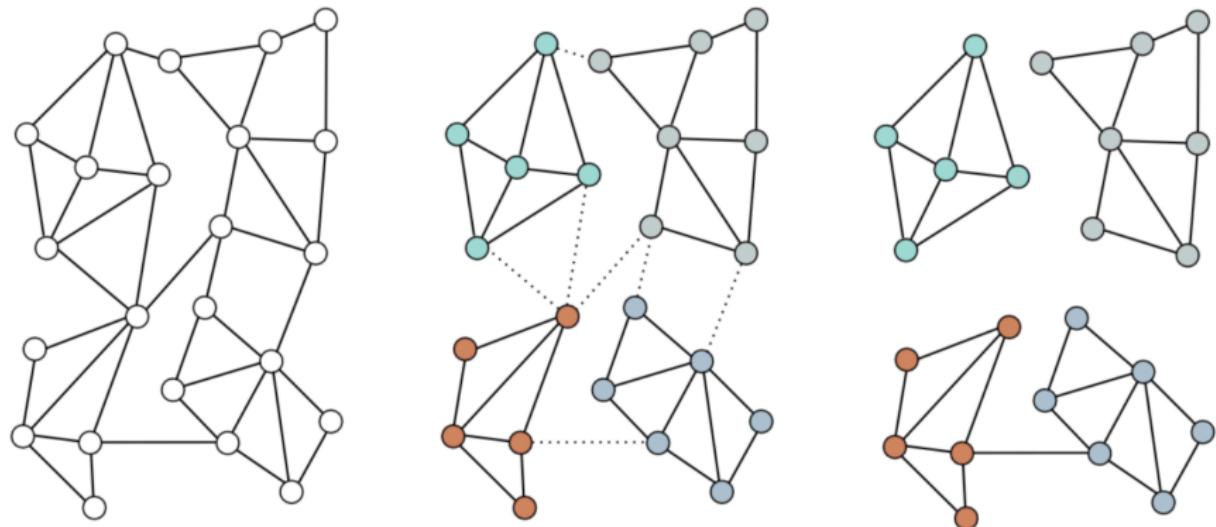
## How it works:

- Cluster the original graph into disjoint subsets of nodes.
- Each subset becomes a smaller subgraph (a "partition") with many internal edges.

## Mini-batch Strategy:

- Treat each partition as a separate training batch.
- Optionally combine multiple partitions in a batch, reintroducing inter-partition edges if needed.

# Graph Partitioning

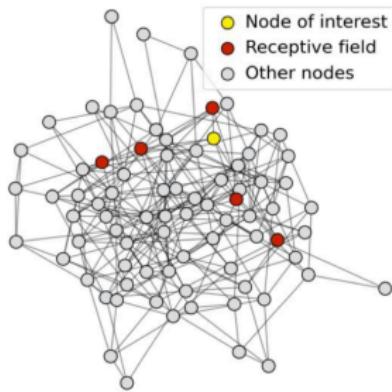


# Over-Smoothing in GNNs

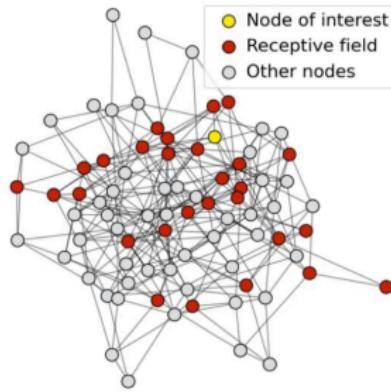
## Problem:

- After multiple layers of message passing, node embeddings tend to become very similar.
- This phenomenon is called **over-smoothing**.
- It limits the expressive power of deep GNNs.

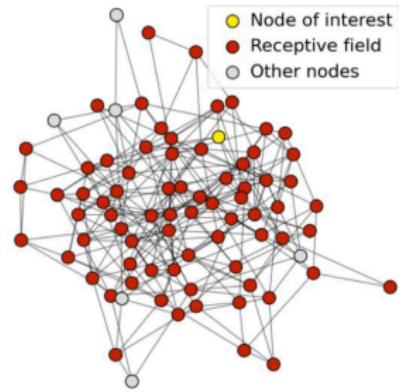
Receptive field for  
**1-layer GNN**



Receptive field for  
**2-layer GNN**



Receptive field for  
**3-layer GNN**



# Over-Smoothing in GNNs

## Why is it bad?

- Node features lose discriminative power.
- The model can no longer distinguish between nodes with different labels or roles.
- This **limits the depth** of the network.

## How to mitigate it:

- **Residual connections:** Preserve original features and stabilize training.
- **Jumping Knowledge connections:** Let the output layer aggregate features from all earlier layers.